

# Deep Reinforcement Learning for Robotic Systems

Mughees Asif

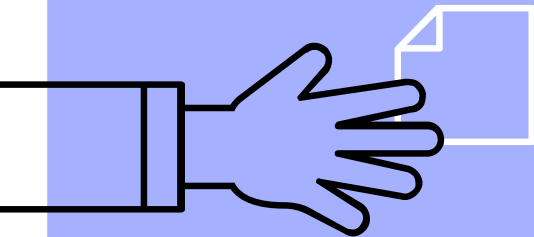
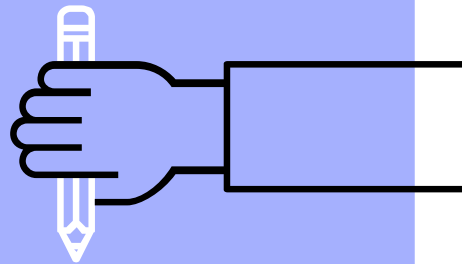
180288337

DEN331: Project Presentation



# Outline

1. Artificial Intelligence (**AI**)
2. Deep Reinforcement Learning (**DRL**)
3. Proximal Policy Optimisation (**PPO**)
4. Current progress and working example
5. Summary



# Artificial Intelligence (AI) Overview

- The capability of a machine to imitate intelligent human behaviour<sup>1</sup>.
- **270%** increase in the use of AI algorithms in the past 4 years<sup>2</sup>.
- Revenue projected to hit **£100 billion** by 2025<sup>2</sup>.
- Use cases involve modelling customer behaviour, streamlining repetitive tasks, and enabling predictive analysis.



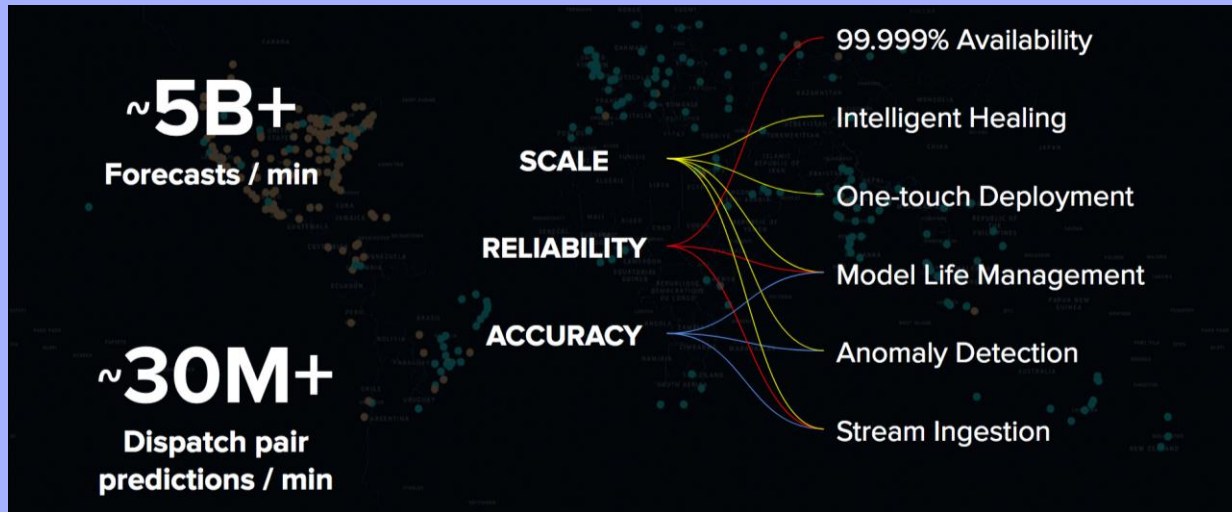
# Artificial Intelligence (AI) Example

- **Miso Robotics:** Flippy.
- A fully autonomous robotic kitchen assistant that uses cloud-based monitoring, thermal imaging and deep learning.
- Improves cooking performance by studying the external environment and food temperature.



# Artificial Intelligence (AI) Example

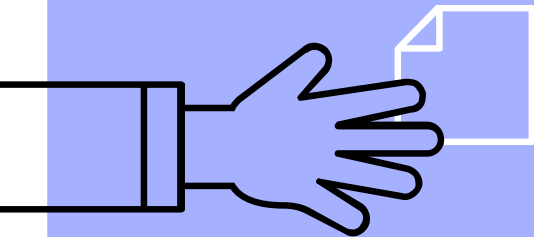
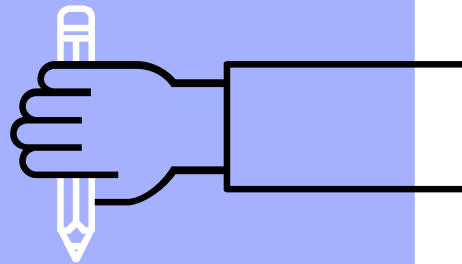
- **Uber:** Michelangelo real-time machine learning system.
- Use cases involve efficient ride-sharing marketplace, identify suspicious or fraudulent accounts, and suggest optimal pickup and drop-off points.



4

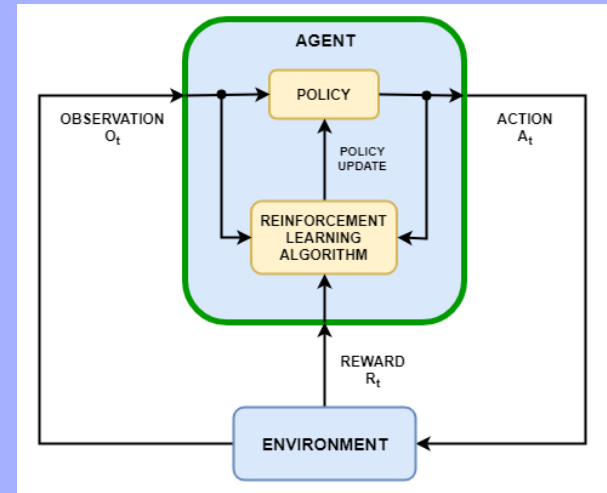
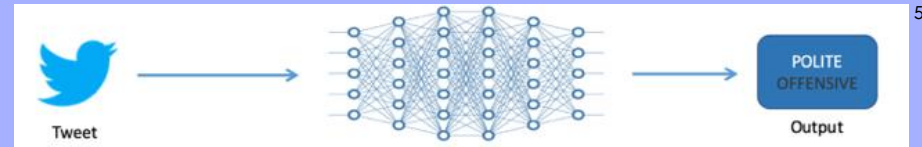
# Outline

1. Artificial Intelligence (**AI**)
2. Deep Reinforcement Learning (**DRL**)
3. Proximal Policy Optimisation (**PPO**)
4. Current progress and working example
5. Summary



# Deep Reinforcement Learning (DRL)

- **Deep Learning:** Abstraction and extraction of pertinent features from a given data set.
- **Reinforcement Learning:** A computational learning model using sequential trial and error.



# Drawbacks

## Neuroplasticity

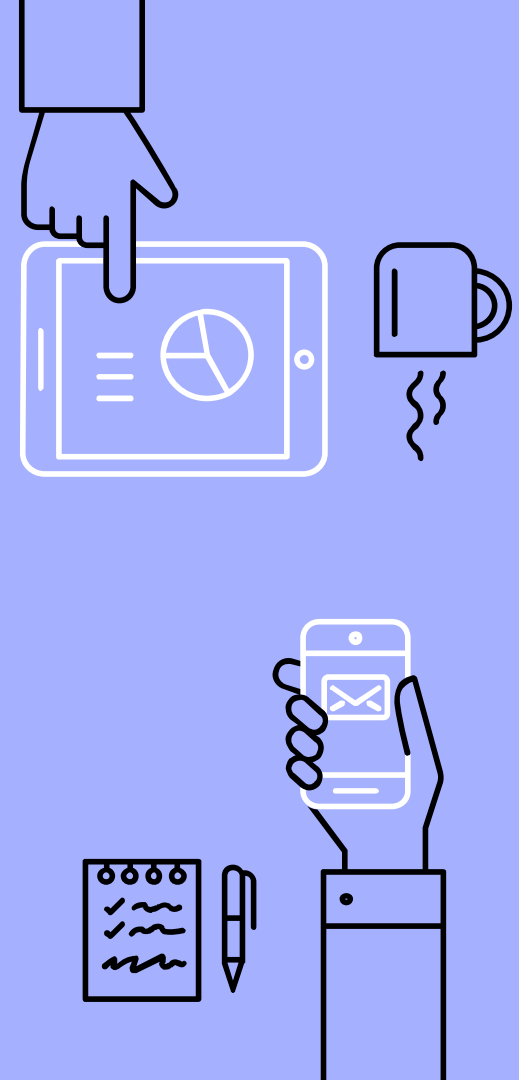
Lack of neuroplasticity decreases the performance, as the complexity is increased.

## Overtraining

Model can predict training examples with very high accuracy, but can not transfer results to new data, leading to poor performance.

## Sample efficiency

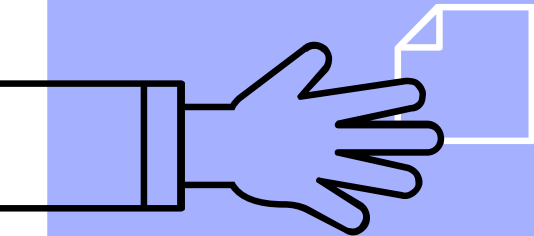
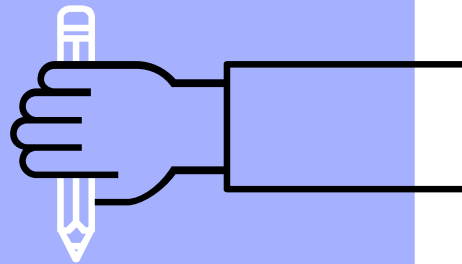
Amount of data needed for training, in order to reach a certain level of performance.





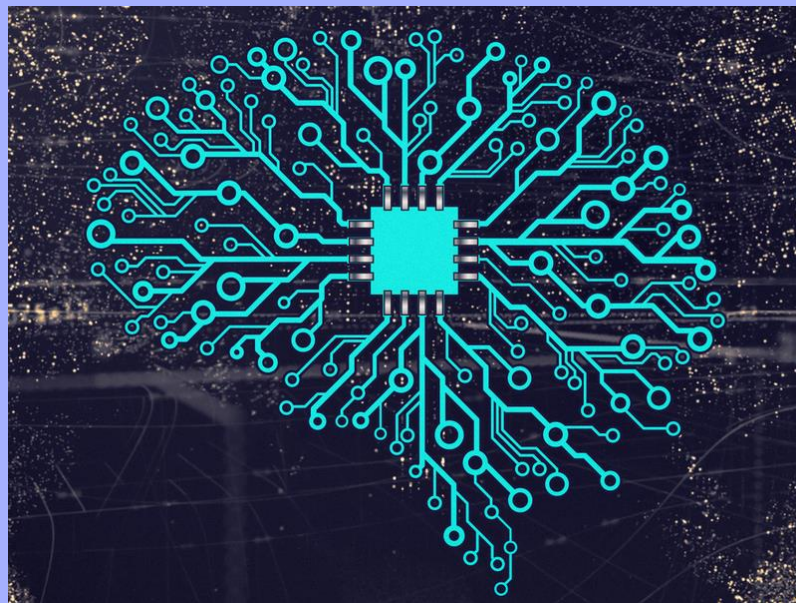
# Outline

1. Artificial Intelligence (**AI**)
2. Deep Reinforcement Learning (**DRL**)
3. Proximal Policy Optimisation (**PPO**)
4. Current progress and working example
5. Summary



# Proximal Policy Optimisation (PPO)

- Developed by **OpenAI** in 2017, to increase accuracy of robotic systems and gaming agents.
- Leverages a **clipped/surrogate function** that maintains stability of learning for the agent, within an **optimal training zone** that has a pre-defined environment range.



7

# Mathematical Modelling

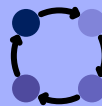
$$\mathbf{L}^{\text{PPO}}(\theta) = \hat{\mathbb{E}}_t [\mathbf{L}^{\text{CLIP}}(\theta) - \mathbf{c}_1 \mathbf{L}^{\text{VF}}(\theta) + \mathbf{c}_2 \mathbf{S}[\pi_\theta](\mathbf{s}_t)]$$

$\mathbf{L}^{\text{CLIP}}(\theta)$



Ensures **reusability** of the policy updates.

$\mathbf{c}_1 \mathbf{L}^{\text{VF}}(\theta)$



Determines **desirability** of the current state.

$\mathbf{c}_2 \mathbf{S}[\pi_\theta](\mathbf{s}_t)$



Ensures **optimum** exploration of an environment by the agent.

# PPO Pseudo-code<sup>8</sup>

## Outer thread

Sporadically collects information to maintain optimum training probability

## Inner thread

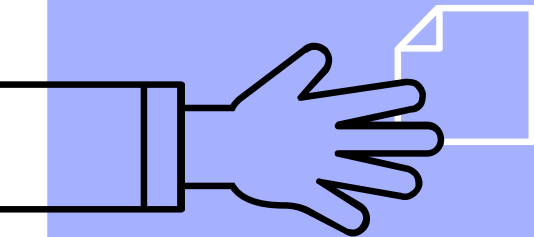
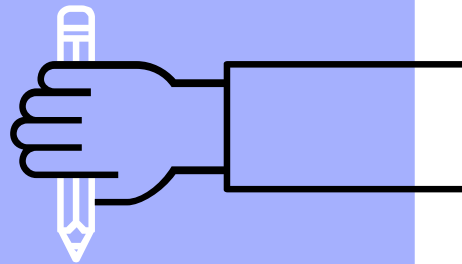
Initiates policy gradient interaction with the simulated environment to calculate and store the desirability of the action

### Algorithm 1 PPO, Actor-Critic Style

```
for iteration=1, 2, ... do
  for actor=1, 2, ..., N do
    Run policy  $\pi_{\theta_{\text{old}}}$  in environment for  $T$  timesteps
    Compute advantage estimates  $\hat{A}_1, \dots, \hat{A}_T$ 
  end for
  Optimize surrogate  $L$  wrt  $\theta$ , with  $K$  epochs and minibatch size  $M \leq NT$ 
   $\theta_{\text{old}} \leftarrow \theta$ 
end for
```


# Outline

1. Artificial Intelligence (**AI**)
2. Deep Reinforcement Learning (**DRL**)
3. Proximal Policy Optimisation (**PPO**)
4. Current progress and working example
5. Summary



# Problem definition and progress

- **Problem:** Use the PPO algorithm to train an agent to balance a double inverted pendulum (DIP).
- **Use cases:** Rocket initialisation trajectory, spacecraft berthing and docking, and stabilisation of human prosthetics.
- **Current:** The equations of motion and computational modelling of the system have been completed.
- **Future:** Implement the PPO algorithm.



10

# Preliminary work

- Used the **SymPy** package to **derive** the dynamics of the double pendulum.

```
qddot_0 = (-4.0*f*cos(2.0*q_2) + 6.0*f + 4.0*qdot_1**2*cos(q_1) + qdot_1**2*cos(q_1 - q_2) - qdot_1**2*cos(q_1 + 2.0*q_2) + 2.0*qdot_1*qdot_2*cos(q_1 - q_2) + qdot_2**2*cos(q_1 - q_2) - 29.43*sin(2.0*q_1) + 9.81*sin(2.0*q_1 + 2.0*q_2))/(3.0*cos(2.0*q_1) - 22.0*cos(2.0*q_2) - cos(2.0*q_1 + 2.0*q_2) + 34.0)
```

```
qddot_1 = (8.0*f*sin(q_1) - 4.0*f*sin(q_1 + 2.0*q_2) + 3.0*qdot_1**2*sin(2.0*q_1) + 23.0*qdot_1**2*sin(q_2) + 22.0*qdot_1**2*sin(2.0*q_2) + qdot_1**2*sin(2.0*q_1 + q_2) + 46.0*qdot_1*qdot_2*sin(q_2) + 2.0*qdot_1*qdot_2*sin(2.0*q_1 + q_2) + 23.0*qdot_2**2*sin(q_2) + qdot_2**2*sin(2.0*q_1 + q_2) - 490.5*cos(q_1) + 215.82*cos(q_1 + 2.0*q_2))/(3.0*cos(2.0*q_1) - 22.0*cos(2.0*q_2) - cos(2.0*q_1 + 2.0*q_2) + 34.0)
```

```
qddot_2 = -((100.0*qdot_1**2*sin(q_2) + 981.0*cos(q_1 + q_2))*(-9.0*(sin(q_1) + 0.3333333333333333*sin(q_1 + q_2))**2 + 28.0*cos(q_2) + 42.0) + 0.5*(200.0*qdot_1*qdot_2*sin(q_2) + 100.0*qdot_2**2*sin(q_2) - 2943.0*cos(q_1) - 981.0*cos(q_1 + q_2))*(25.0*cos(q_2) + 3.0*cos(2.0*q_1 + q_2) + cos(2.0*q_1 + 2.0*q_2) + 13.0) + 50.0*(2.0*sin(q_1) + 3.0*sin(q_1 - q_2) - 2.0*sin(q_1 + q_2) - sin(q_1 + 2.0*q_2))*(2.0*f + 3.0*qdot_1**2*cos(q_1) + qdot_1**2*cos(q_1 + q_2) + 2.0*qdot_1*qdot_2*cos(q_1 + q_2) + qdot_2**2*cos(q_1 + q_2)))/(75.0*cos(2.0*q_1) - 550.0*cos(2.0*q_2) - 25.0*cos(2.0*q_1 + 2.0*q_2) + 850.0)
```

- Used the **do-mpc** package to **solve** the Euler-Lagrangian system dynamics of the double pendulum.

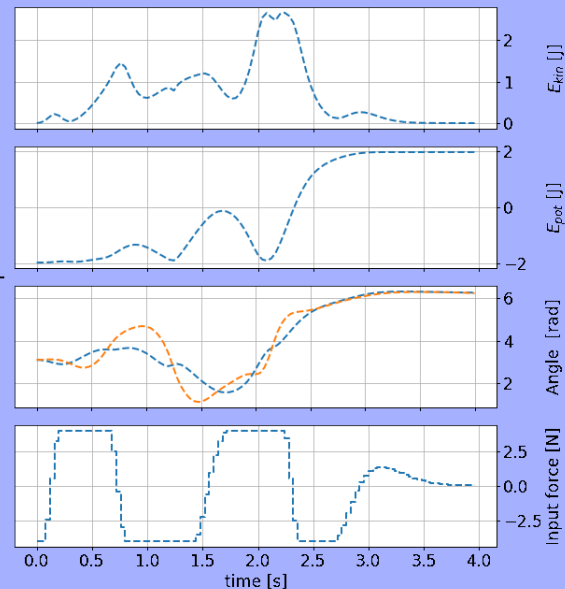

$$\begin{aligned}h_1\ddot{x} + h_2\ddot{\theta}_1 \cos(\theta_1) + h_3\ddot{\theta}_2 \cos(\theta_2) &= (h_2\dot{\theta}_1^2 \sin(\theta_1) + h_3\dot{\theta}_2^2 \sin(\theta_2) + u) \\h_2 \cos(\theta_1)\ddot{x} + h_4\ddot{\theta}_1 + h_5 \cos(\theta_1 - \theta_2)\ddot{\theta}_2 &= (h_7 \sin(\theta_1) - h_5\dot{\theta}_2^2 \sin(\theta_1 - \theta_2)) \\h_3 \cos(\theta_2)\ddot{x} + h_5 \cos(\theta_1 - \theta_2)\ddot{\theta}_1 + h_6\ddot{\theta}_2 &= (h_5\dot{\theta}_1^2 \sin(\theta_1 - \theta_2) + h_8 \sin(\theta_2))\end{aligned}$$

<https://github.com/mughees-asif/dip#third-year-thesis-project>

# Working example #1

*# model parameters*

```
m0 = 0.6 # kg, mass of the cart  
m1 = 0.2 # kg_1, mass of the first rod  
m2 = 0.2 # kg_2, mass of the second rod  
L1 = 0.5 # m_1, length of the first rod  
L2 = 0.5 # m_2, length of the second rod  
g = 9.81 # m/s^2, Gravity
```

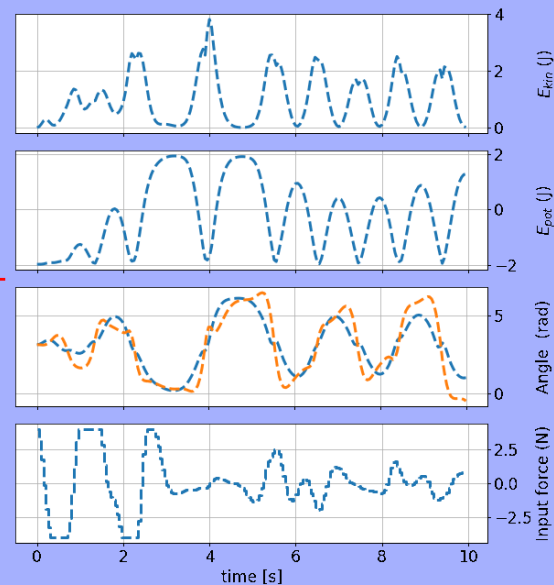
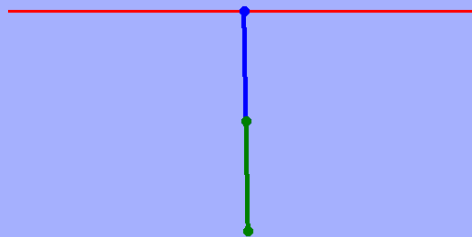


<https://github.com/mughees-asif/dip#third-year-thesis-project>



# Working example #2

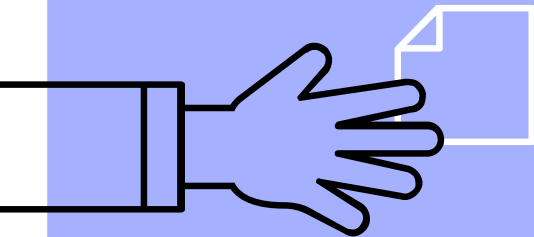
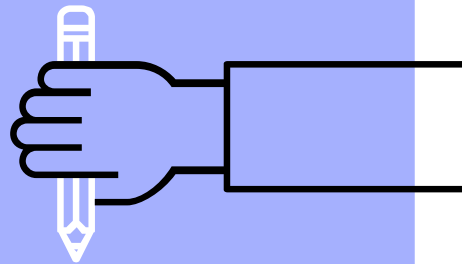
```
# model parameters
m0 = 3 # kg, mass of the cart
m1 = 1 # kg_1, mass of the first rod
m2 = 1 # kg_2, mass of the second rod
L1 = 0.5 # m_1, length of the first rod
L2 = 0.5 # m_2, length of the second rod
g = 9.81 # m/s^2, Gravity
```



<https://github.com/mughees-asif/dip#third-year-thesis-project>

# Outline

1. Artificial Intelligence (**AI**)
2. Deep Reinforcement Learning (**DRL**)
3. Proximal Policy Optimisation (**PPO**)
4. Working example with context to the final-year project
5. Summary



# Summary

## Artificial Intelligence

The ability of a machine to perceive its environment, whilst *reacting* in a way that successfully maximises a user-defined output.

## Deep Reinforcement Learning

A synchronous combination of leveraging the advantages of using deep learning (DL) with reinforcement learning (RL).

## Proximal Policy Optimisation

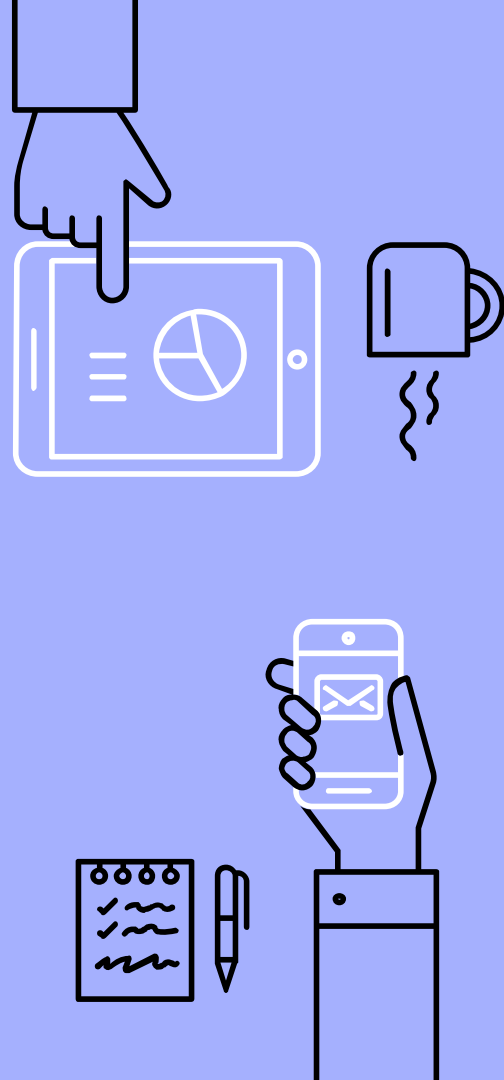
A set of deep reinforcement learning algorithms that use clipped/surrogate objectives to ensure the agent stays within an optimal training zone.

## Current progress

Modelling of the system has been completed and the initial set of findings have been displayed.

## Future steps

Develop the PPO algorithm that is applicable to this system and train an agent to execute the task of balancing a double inverted pendulum.



# THANKS!

## Any questions?



# Citations

- <sup>1</sup> Merriam-Webster. n.d. *Definition Of Artificial Intelligence*. [online] Available at: <<https://www.merriam-webster.com/dictionary/artificial%20intelligence>> [Accessed 24 November 2020].
- <sup>2</sup> Lin, Y. *10 Artificial Intelligence Statistics You Need to Know in 2020* . [Infographic]. Nov. 2020. Available at: <https://www.oberlo.co.uk/blog/artificial-intelligence-statistics>.
- <sup>3</sup> Gadgetify. n.d. *Flippy Burger Flipping Robotic Kitchen Assistant*. [online] Available at: <<https://www.gadgetify.com/flippy-burger-robot/>> [Accessed 24 November 2020].
- <sup>4</sup> Turakhia, C., 2017. *Engineering More Reliable Transportation With Machine Learning And AI At Uber*. [online] Uber Engineering Blog. Available at: <<https://eng.uber.com/machine-learning/>> [Accessed 24 November 2020].
- <sup>5</sup> Paradzhanyan, L., Gunter, A. and Brotsos, J., 2020. *Demystifying Deep Learning And Artificial Intelligence* . [online] The New Stack. Available at: <<https://thenewstack.io/demystifying-deep-learning-and-artificial-intelligence/>> [Accessed 25 November 2020].
- <sup>6</sup> Mathworks. n.d. *Reinforcement Learning Agents*. [online] Available at: <<https://uk.mathworks.com/help/reinforcement-learning/ug/create-agents-for-reinforcement-learning.html>> [Accessed 25 November 2020].
- <sup>7</sup> Karidis, V., 2020. *Proximal Policy Optimization: AI'S Best Way To Transfer Knowledge?* [online] Ocular Point. Available at: <<https://www.ocularpoint.com/post/proximal-policy-optimization-ai-s-best-way-to-transfer-knowledge>> [Accessed 26 November 2020].
- <sup>8</sup> J. Schulman, F. Wolski, P. Dhariwal, A. Radford, and O. Klimov. *Proximal policy optimization algorithms*. CoRR, vol. abs/1707.06347, 2017. [Online]. Available: <http://arxiv.org/abs/1707.06347>.
- <sup>9</sup> Wang C, Zhang Q, Tian Q, Li S, Wang X, Lane D, Petillot Y, Wang S. *Learning Mobile Manipulation through Deep Reinforcement Learning*. Sensors. 2020; 20(3):939.
- <sup>10</sup> do-mpc. n.d. *Double Inverted Pendulum — Do-Mpc 4.0.0 Documentation*. [online] Available at: <[https://www.do-mpc.com/en/latest/example\\_gallery/DIP.html](https://www.do-mpc.com/en/latest/example_gallery/DIP.html)> [Accessed 27 November 2020].